

Quiz Computer Vision

Kelompok 10

Anggota Kelompok:

Nama : Rio Pramana, Debora

NIM : 2440016804, 2440087214

Kelas : LA01

Import libraries

```
In [ ]: #import libraries
import cv2
from matplotlib import pyplot as plt
import numpy as np
from os import listdir
from scipy.cluster.vq import *
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
```

Load dataset

Our device was not able to handle the processing of all 37500 images of the dataset. Therefore, we had to reduce the number of data to 10% of the original training dataset (2500 training data). We only took 10 sample image from the test dataset (from 12500 images) because all of the test dataset are unlabeled, so we only took several images just to show how well the model/algorithm could perform to predict whether an image is a cat or a dog

```
In [2]: #Read dataset
TRAIN_PATH = '../Images/train/reduced_train'
TEST_PATH = '../Images/reduced_test'

train_filename_list = listdir(TRAIN_PATH)
```

We also resized each image to 128 x 128 to make the image processing easier

```
In [3]: def get_dataset(path, is_train=1):
        filename_list = listdir(path)
        images = []
        classes = []

        for filename in filename_list:
            image = cv2.resize(cv2.imread(f"{path}/{filename}"), (128,128))
            images.append(image)
            if is_train == 1:
                img_class = filename.split('.')[0]
                if img_class == "dog":
                    id = 1
                else:
                    id = 0
                classes.append(id)

        return images, classes
```

Get both dataset:

```
In [4]: train_images, train_classes = get_dataset(TRAIN_PATH)
        test_images, _ = get_dataset(TEST_PATH, 0)
```

Collect all the descriptors from the images

First, we need to create/instantiate a SIFT object (We are using SIFT class from opencv that is used to extract keypoints and computing descriptors using SIFT algorithm proposed by D. Lowe)

Then, we use that object to only get the descriptors from each image

Below is the function to get descriptors from each image in a list of images (for example: list of train images)

```
In [5]: #Extract features
sift = cv2.SIFT_create()

def get_descriptors(img_list):
    descriptors = []

    for image in img_list:
        _, desc = sift.detectAndCompute(image, None)
        descriptors.append(desc)

    return descriptors
```

Get descriptors from training & test dataset:

```
In [6]: train_desc = get_descriptors(train_images)
        test_desc = get_descriptors(test_images)
```

```
In [7]: print(len(train_desc))
```

2500

Before moving on, we need to apply vertical stack to stack the sequence of arrays vertically into a single array. Therefore, each training image's descriptors will be collected/put into a single list (not lists in a list) which will act as the **bag of descriptors**

```
In [8]: desc_bag = train_desc[0]
        for descriptor in train_desc[1:]:
            desc_bag = np.vstack((desc_bag, descriptor))

        desc_bag = np.float32(desc_bag)
```

```
In [9]: print(len(desc_bag))
```

323773

Cluster descriptors into k number of clusters

For k, we put it the same as the number of categories/classes in this dataset, which is 2

We used the default number of iterations (20)

```
In [10]: centroids, _ = kmeans(desc_bag,2,20)
```

For every image in the dataset, counts how frequently the different features in the dictionary occurs

For each SIFT descriptor in the image, we find the closest cluster center (centroids) and increment its corresponding count in the vector v by 1 using vector quantization from scipy

```
In [11]: def get_img_features(img_list, centroids, descriptors):
        image_features = np.zeros((len(img_list), len(centroids)), "float32")
        for i in range(len(img_list)):
            words, dist = vq(descriptors[i], centroids)
            for w in words:
                image_features[i][w] +=1

        return image_features
```

```
In [12]: train_features = get_img_features(train_images, centroids, train_desc)
        test_features = get_img_features(test_images, centroids, test_desc)
```

Normalize the histogram

To normalize the features, we use StandardScaler()

```
In [13]: scaler = StandardScaler()
        scaler = scaler.fit(train_features)
        train_features = scaler.transform(train_features)
        test_features = scaler.transform(test_features)
```

Create k-NN model

To predict test images, we use k-NN classifier with n_neighbors = 1 because it gave the best results from our experiments (we tried the value 1 - 6 and the best performing one is with n_neighbors = 1)

```
In [36]: knn = KNeighborsClassifier(n_neighbors = 1)
        knn.fit(train_features, np.array(train_classes))
```

```
Out[36]: KNeighborsClassifier(n_neighbors=1)
```

Evaluate Performance

Because the test dataset are all unlabeled, we first tried to predict the train dataset

```
In [37]: train_pred = knn.predict(train_features)
        print(classification_report(train_classes, train_pred, target_names = ['cat', 'dog']))
```

	precision	recall	f1-score	support
cat	0.88	0.88	0.88	1250
dog	0.88	0.88	0.88	1250
accuracy			0.88	2500
macro avg	0.88	0.88	0.88	2500
weighted avg	0.88	0.88	0.88	2500

It managed to achieve the same performance (88%) on all aspects

Next, we tried to predict the 10 test images (we picked the first 5 dog images and 5 cat images from the test dataset)

```
In [38]: test_pred = knn.predict(test_features)
```

Show the prediction:

```
In [39]: plt.figure(figsize=(12,6))
        for id, (result, image) in enumerate(zip(test_pred, test_images)):
            plt.subplot(2, 5, id+1)
            if test_pred[id] == 1:
                plt.title("dog")
            else:
                plt.title("cat")
            plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
            plt.xticks([])
            plt.yticks([])
        plt.show()
```



The model was able to achieve a good performance (accurately predict 8/10 test images). However, the model performance could probably be even higher if it was able to use the whole dataset (37500 images) without having to resize the images (to 128 x 128) because it might also affect the images' descriptors collection